

Reducing Leakage Power in Cache Peripheral Circuits of Embedded Processors

HOUMAN HOMAYOUN

COMPUTER SCIENCE AND ENGINEERING DEPARTMENT
UNIVERSITY OF CALIFORNIA SAN DIEGO

ALEX VEIDENBAUM

COMPUTER SCIENCE DEPARTMENT
UNIVERSITY OF CALIFORNIA IRVINE

AVESTA SASAN

BROADCOM CORPORATION, OFFIC OF CTO

DEAN TULLSEN

COMPUTER SCIENCE AND ENGINEERING DEPARTMENT
UNIVERSITY OF CALIFORNIA SAN DIEGO

Recent studies have shown that peripheral circuits, including decoders, wordline drivers, input and output drivers, contribute a large fraction of the overall cache leakage. In addition, as technology migrates to smaller geometries, leakage contribution to total power consumption increases faster than dynamic power, indicating that leakage will be a major contributor to overall power consumption. This paper proposes a combination of circuit and architectural techniques to maximize leakage power reduction in embedded processors on-chip caches by targeting leakage in cache peripheral circuits. At the circuit level, we propose a novel design with multiple sleep modes for cache peripherals. Each mode represents a trade-off between leakage reduction and wakeup delay. Architectural control is proposed to decide when and how to use these different low-leakage modes by utilizing cache miss information to guide its action. The control is based on simple state machines that do not impact area or power consumption and can thus be used even in resource-constrained processors. Experimental results indicate that the proposed techniques can keep the data cache peripherals in one of the low-power modes for more than 85% of total execution time. The instruction cache is in one of the low power modes for an average of 50% of the time. This translates to an average leakage power reduction of 50% for DL1 and 30% for IL1 in 65nm technology. The energy-delay product is reduced, on average, by up to 21%. The overall processor power is reduced by up to 18% and by an average of 11% to 14%.

Categories and Subject Descriptors B.3.2 [MEMORY STRUCTURES], Design Styles: Cache memories; C.1.1 [PROCESSOR ARCHITECTURES], Single Data Stream Architectures: Pipeline processors Systems

General Terms: Design, Performance

Additional Key Words and Phrases: Cache, Leakage Power, Peripheral Circuits, Multiple Sleep Mode, Embedded Processor

1. INTRODUCTION

Static or leakage energy consumption has been growing in both embedded and high-performance processors as transistor geometries shrink. Cache and TLB RAM structures account for a large fraction of processor power consumption [Rusu et al. 2007], [Montanaro et al. 1996], and especially for leakage power. A number of process and circuit techniques have been proposed to significantly reduce leakage of the memory cell array making SRAM peripheral circuits the main sources of leakage. Recent results have shown a large amount of leakage in the peripheral SRAM circuits, such as decoders, word-line and output drivers [Homayoun et al. 2008], [Nakagome et al. 2003] [Takeyama et al. 2006], [Gerosa et al. 2008], [Nii et al. 2004] [Mamidipaka et al. 2004]. For instance, a wordline driver drives its signal to a large number of memory cells. To drive a

high capacitive load, a chain of tapered inverter buffers is used, typically with three to five levels. We compare the leakage power consumption of a 65nm SRAM6T memory cell¹ with inverters of different size. The results are shown in Figure 1(a). It

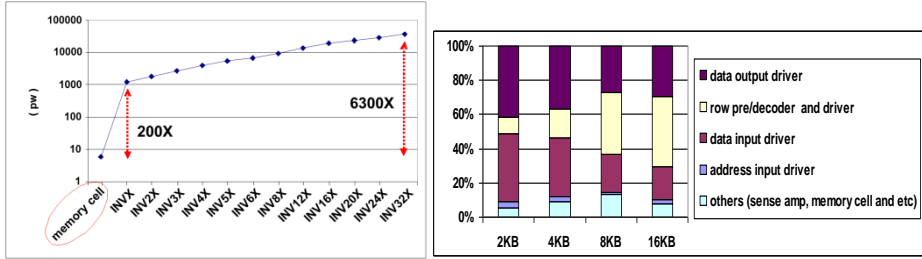


Figure 1. (a) Leakage power dissipation of one SRAM6T memory cell compared with different size inverter buffers (INVX is the smallest inverter buffer with drive strength of 1) (b) Leakage power components for different cache size

shows that the leakage power of a standard memory cell is significantly lower than the inverter buffer leakage and that the inverter leakage grows exponentially with its size. Let us assume that a driver has to drive 256 one-bit memory cells. This will require three stages of inverter buffers, with buffer size increasing by a factor of e in each additional stage. The combined leakage power of these three drivers is 12 times larger than the leakage of the 256 memory cells. In addition to the wordline driver, one has to consider leakage in data input and output drivers, which are also high.

Two main reasons explain this difference in leakage:

- Memory cells are designed with minimum-sized transistors mainly for area considerations. Unlike memory cells, peripheral circuits use larger, faster and accordingly, more leaky transistors so as to satisfy timing requirements [Nakagome et al. 2003], [Mamidipaka et al. 2004].
- Memory cells use high threshold voltage transistors, which have a significantly lower leakage reduction compared with typical threshold voltage transistors used in peripheral circuits [Takeyama et al. 2006], [Mamidipaka et al. 2004].

In summary, SRAM memory cells are optimized for low leakage current and area without a significant impact on the cell performance [Nakagome et al. 2003], [Takeyama et al. 2006], [Grossara et al. 2006], [Nii et al. 2004]. In addition, circuit techniques such as gated- V_{dd} and drowsy cache can be applied to further reduce the memory cell leakage and further widen the gap between the cell array and peripheral circuit leakage power dissipation.

A similar result was obtained using CACTI 5.1 [Thoziyoor et al. 2008]. CACTI uses characteristics of transistors modeled by the ITRS [Semiconductor Industries Association, 2005]. It includes data for two device types that the ITRS defines - High Performance (HP) and Low Standby Power (LSTP). The HP transistors are fast transistors with short gate lengths, thin gate oxides, low V_{th} , and low V_{dd} . The LSTP transistors, on the other hand, have longer gate lengths, thicker gate oxides, higher V_{th} , and higher V_{dd} . As explained in [Nakagome et al. 2003], HP transistors are used in the peripheral circuitry while LSTP transistors are used in the memory cell array. While it is

¹ similar results were obtained for TSMC, TOSHIBA, IBM, UMC and CHARTERED foundries using their libraries and evaluating leakage with SPICE.

possible to use LSTP transistors in peripheral circuits for reducing leakage, the impact on memory access time would be significant; for instance, it would result in an increase from 1.3ns to 5.7ns in access time of a 16KB L1 cache. Figure 1(b) shows leakage components for different size SRAMs in 65nm technology (based on CACTI 5.1 [Thoziyoor et al. 2008]), with peripheral circuits – data drivers, address driver, decoder and wordline drivers – accounting for over 80% of overall SRAM leakage. Thus, approaches that concentrate on cell leakage power alone are insufficient and it is very important to address leakage in the peripheral circuits.

Since peripheral circuits use very large transistors to drive high loads to meet the memory timing constraints, applying traditional leakage reduction techniques such as “fine grained” sleep transistor insertion [Calhoun et al. 2003], [Khandelwal et al. 2004] in these circuits could introduce significant area and delay overhead. This is mainly due to the impact of leakage reduction techniques on peripheral circuits’ rise time, fall time, propagation delay and area overhead [Homayoun et al. 2008]. These delays not only can significantly increase the cache access time, but also require significant redesign and verification efforts to avoid impacting memory functionality [Homayoun et al. 2008]. In addition, peripheral circuits drive specific voltage levels and cannot be simply power-gated. For instance, a word-line driver has to drive a V_{low} voltage for all wordlines not being accessed.

Therefore, the focus of this paper is on reduction of leakage-related power dissipation in the on-chip SRAMs, targeting their peripheral circuits. The paper explores an *integrated* circuit and architectural approach to reduce leakage in the cache peripherals. It extends the work presented in [Homayoun et al. 2008], highlighting the importance of leakage power reduction in cache peripherals, since they have the highest leakage energy in the processors. There are a large variety of embedded processors, from single-issue, in-order processors with one level of cache to multiple-issue, out-of-order processors with two levels of cache. For this work we define the former as low-end embedded processors, while a single-issue, in-order processor with two cache levels is defined as high-end. The two types of embedded processors also differ in their clock speed. Our goal is to explore leakage reduction in different levels of cache hierarchy. The same techniques are applicable to out-of-order embedded processors that also offer other opportunities for cache leakage reduction, but this type of processor is beyond the scope of this paper.

At the circuit level, we utilize the recently proposed *zig-zag share* circuit [Homayoun et al. 2008] to reduce the sub-threshold leakage in peripheral circuits of L1 caches (IL1 and DL1). The *zig-zag horizontal and vertical share* technique was shown to be very effective in reducing leakage of SRAM peripherals. The results in [Homayoun et al. 2008] showed leakage reduction by up to 100X in deeply pipelined SRAM peripheral circuits, with only a minimal area overhead and small additional delay.

The wakeup latency of *zig-zag share* could be large, especially in large SRAMs, as shown in [Homayoun et al. 2008]. To deal with this problem, this paper shows that by increasing the gate bias voltage of the NMOS footer sleep transistor in the *zig-zag share* circuit (and decreasing it for the PMOS header transistor) one can trade leakage reduction for wakeup delay. Thus we propose to use several low-leakage modes with different wakeup times to better control the cache leakage. For instance, one can have a low-leakage mode for a DL1 cache with a one-cycle wakeup but it would reduce leakage by only 40%. Alternatively, one can define a mode with a 4-cycle wakeup that saves 90% of leakage. These modes differ only in how they bias sleep transistors and thus can be

dynamically switched during execution with almost no delay. We also propose a low-power design for the bias generator circuit required for the multiple sleep modes operation, shown to be fairly robust against six-sigma process, voltage and temperature variations. In addition, we explore the design space of sleep transistor insertion in SRAM peripheral circuitry and show the effect of sleep transistor size, its gate bias and the number of horizontal and vertical levels of sharing on the leakage power savings, instability, dynamic power, propagation delay, wakeup delay, rise time and fall time of the SRAM peripheral circuits.

The next question is when and how to use these different low-leakage modes for L1 caches. Note that this approach can also be applied to L2 caches, but this is beyond the scope of this paper. We propose architectural control of low-leakage modes in L1 caches. The policy uses cache miss information to determine the mode. The action depends in part on the cache organization and in part on the ability to hide the wakeup delay. In all cases, control is based on simple state machines that do not impact area or power consumption and can thus be used even in low-end processors. Hiding one to four cycles of wakeup latency in a short pipeline typical of embedded processors is difficult to accomplish in a uniform way, therefore we propose different methods for each delay and low-leakage mode. For instance, one cycle of delay in DL1 access can be completely hidden by starting cache wakeup as soon as instruction type is known in decode. Thus we can actually keep the L1 cache peripherals in a basic low power mode with a one-cycle wakeup as default. The most efficient low-energy mode with a four cycle wakeup delay can be used during cache miss service. Other low-power modes can be used when an L2 cache is present.

The energy savings and performance of various cache configurations for embedded processors are evaluated in this paper using the proposed circuit and architectural techniques.

It is shown that our techniques can keep the data cache peripherals in one of the low-power modes for more than 85% of total execution time, on average. The instruction cache is in one of the low power modes for an average of 50% of the time. This translates to an average leakage power reduction of 50% for DL1 and 30% for IL1 in 65nm technology. The energy-delay product is reduced, on average, by up to 21%. The overall processor power is reduced by up to 18.6% and by an average of 11% to 14%.

2. RELATED WORK

A number of techniques have been proposed for reducing leakage power at technology, circuit, architecture and compiler/OS levels.

2.1 Circuit-Level Leakage Control

Several circuit techniques have been proposed to reduce the leakage power in SRAM memories. These techniques mainly target the SRAM memory cell leakage.

The primary technique is voltage scaling which, due to short-channel effects in deep submicron processes reduces the leakage current significantly [Flautner et al. 2002]. Another technique is Gated-Vdd which turns off the supply voltage of memory cells by using a sleep transistor and eliminating the leakage almost completely [Powell et al. 2000]. However, it doesn't retain the state of the memory cells. The third technique, ABB-MTCMOS, increases the threshold voltage of an SRAM cell dynamically through controlling its body voltage [Nii et al. 1998]. The performance and area overhead of

applying this technique makes it inefficient. Device scaling leads to threshold voltage fluctuation, which makes the cell bias control difficult to achieve. In response, [Takeyama et al. 2006] propose a Replica Cell Biasing scheme in which the cell bias is not affected by V_{dd} and V_{th} of peripheral transistors.

[Agarwal et al. 2002] and [Kim et al. 2005] propose a forward body biasing scheme (FBB) in which the leakage power is suppressed in the unselected memory cells of cache by utilizing super V_t devices.

In addition to these four major techniques applied to SRAM memories, there are also leakage reduction techniques in literature which concentrate on generic logic circuits. Examples are sleepy stack [Park et al. 2006], sleepy keeper [Kim et al. 2006] and zig-zag super cut-off CMOS (ZSCCMOS) techniques [Min et al. 2003], [Horiguchi et al. 1993]. ZSCCMOS reduces the wakeup overhead associated with the Gated- V_{dd} technique by inserting the sleep transistors in a zig-zag fashion. Sleepy stack divides the existing transistors into two half-size and then inserts sleep transistor to further reduce leakage. This approach was shown to be area-inefficient as it comes with 50 to 120% area overhead. Sleepy keeper is a variation of Gated- V_{dd} approach which can save logic state during sleep mode. The drawback of this approach is significant additional dynamic power overhead compare to the base circuit.

Optimal sizing of sleep transistors to minimize the impact of sleep transistor insertion on the circuit delay has been researched extensively [Calhoun et al. 2003], [Mutoh et al. 1999], [Khandelwal et al. 2004], [Ramalingam et al. 2005]. [Mutoh et al. 1999] used the average current consumption of logic gates to find the size of sleep transistor for satisfying circuit speed. Their proposed technique is based on the assumption that the circuit speed is weakly dependent on the circuit operating pattern for large enough sleep transistor size.

2.2 Architectural Techniques

A number of architecturally driven cache leakage reduction techniques have been proposed. Powell et al applied gated- V_{dd} approach to gate the power supply for cache lines that are not likely to be accessed [Powell et al. 2000]. Kaxiras et al. propose a cache decay technique which reduces cache leakage by turning off cache lines not likely to be reused [Kaxiras et al. 2001]. Flautner et al. proposed a drowsy cache which reduces the supply voltage of the L1 cache line instead of gating it off completely [Flautner et al. 2002]. The advantage of this technique is that it preserves the cache line information but introduces a delay in accessing drowsy lines. Nicolaescu et al [Nicolaescu et al. 2006] proposed a combination of way cache technique and fast speculative address generation to apply the drowsy cache line technique to reduce both the L1 cache dynamic and leakage power. [Zhang et al. 2002] proposed a compiler approach to turn off the cache lines for a region of code that would not be accessed for a long period of time [Meng et al. 2005]. Meng et al. presented a perfecting scheme which combines the drowsy caches and the Gated- V_{dd} techniques to optimize cache leakage reduction [Meng et al. 2005]. Ku et al. [Ku et al. 2006] exploit several power density minimization techniques to reduce temperature and further leakage in highly-associative

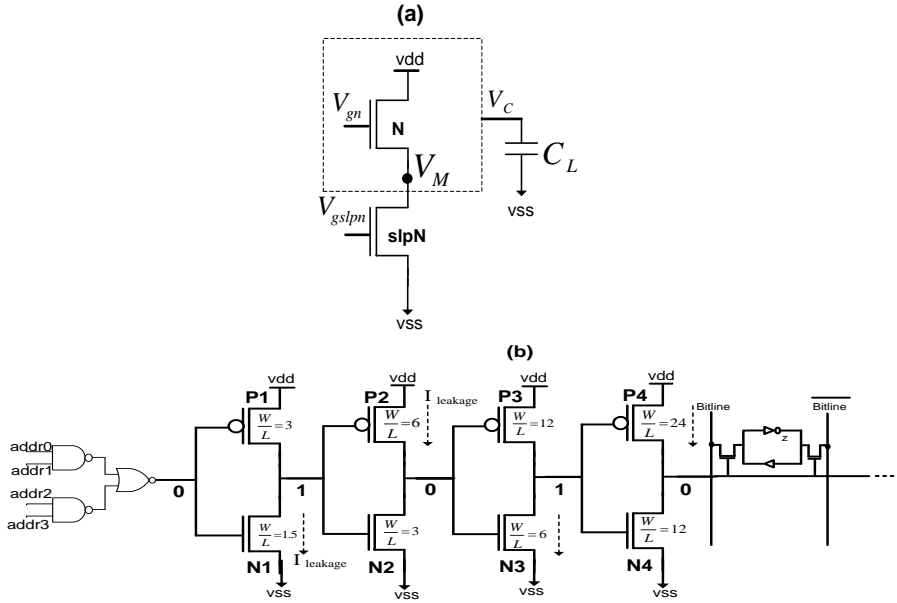


Figure 2. (a) Stacking sleep transistor to reduce leakage (b) Leakage in the wordline driver

on-chip caches in embedded processors. Due to the positive feedback relation of temperature and leakage, they show on-chip cache leakage reduces significantly.

In a recent work, Chun et al. proposed a novel approach, which utilizes branch prediction assist to selectively wake up only the needed instruction cache line for an embedded processor [Chun et al. 2008].

All research mentioned above primarily targets the leakage in the SRAM cells of a cache. Given the results in Figure 1 and in [Homayoun et al. 2008], [Gerosa et al. 2008], [Takeyama et al. 2006], [Nii et al. 2004] and [Mamidipaka et al. 2004], peripheral circuits are equally if not more important to address in a cache.

3. SLEEP TRANSISTOR STACKING

Stacking sleep transistors have been proposed to reduce sub-threshold (I_{Dsub}) or weak inversion current [Kao et al. 1998]. I_{Dsub} is an inverse exponential function of threshold voltage (V_T). Threshold voltage is a function of Source to Bulk Voltage. An effective way to reduce the leakage of a transistor is by increasing its source voltage (for an NMOS increasing V_{SB} , the source to bulk voltage) [Kao et al. 1998], [Rabaey et al. 2003].

Stacking a sleep transistor (footer NMOS or header PMOS transistor) as shown in Figure 2(a) could deliver this effect. In this figure by stacking transistor N with slpN source to body voltage (V_M) of transistor N increases. When both transistors are off an increase in V_M increases the V_T of the transistor N and therefore reduces sub-threshold leakage current [Kao et al. 1998]. Size (W/L) and bias (V_{gslpn}) voltage of the stacked sleep transistor determines the V_M [Agarwal et al. 2006], [Kao et al. 1998]. Reducing the sleep transistor bias reduces the leakage while increasing the circuit wakeup period which is the time to pull the V_M down to ground. Thus there is a trade-off between the amount of leakage saved and the wakeup overhead [Agarwal et al. 2006].

A wordline driver shown in Figure 2(b) increases the gate voltage of the access transistors of all cells connected to the selected wordline. The number and size of inverters in the chain are chosen to meet the timing requirements for charging or discharging the wordline (usually the number of inverter stages varies from two to five. For this particular example we assumed four stages. Based on our experimental results the relative leakage power consumption does not change noticeably for our case, compared to a design with fewer or more inverter stages). The size of inverters in the chain decreases from decoder side to the wordline to increase the effective fan-out. The inverter chain has to drive a logic value 0 to the pass transistors when a memory row is not selected. Thus the driver cannot be simply shut down when idle. Transistors N1, N3 and P2, P4 are in the off state and thus they are leaking.

Stacking header and footer sleep transistors with all NMOS and PMOS transistors in the chain reduces their leakage; however, aside from the area overhead, it increases the propagation delay of the inverters in the driver chain followed by an increase in the rise/fall time of the wordline. Rise and fall time of an inverter output is proportional to $R_{peq} * C_L$ and $R_{neq} * C_L$, respectively, where R_{peq} is the equivalent resistance of the PMOS transistor, R_{neq} is the equivalent resistance of the NMOS transistor, and C_L is the equivalent wordline output capacitive load [Rabaey et al. 2003]. Inserting sleep transistors increases R_{neq} , R_{peq} and thus the rise time and fall time of the wordline driver as well as its propagation delay [Homayoun et al. 2008], [Rabaey et al. 2003]. Increase in rise and fall times and propagation delay is not desirable, as the memory functionality and access time are negatively affected [Amrutur et al. 2000], [Mai et al. 1998], [Margala et al. 1999]. Any increase in the fall time of wordline driver results in an increase in pass transistor active period during a read operation. This results in the bit-line over-discharge and the memory content over-charge during the read operation. Such over-discharge not only increases the dynamic power dissipation of bit-lines, more importantly, it can cause a memory cell content to enter a metastable state if the over-discharge period is large [Amrutur et al. 2000], [Mai et al. 1998], [Margala et al. 1999]. In addition, such increase in pass transistors active period requires re-timing of the sense amplifier active period.

Note that it is typical in low power SRAM memories to use a self-timed clocked sense amplifier to limit sense power and track the bit-line delay to setup the sense amplifier in the high gain region for sensing right before the pass transistors are turned off. Such a self-timed clocked sense amplifier has to cope with any increase in the pass transistor M1 active period [Amrutur et al. 2000], [Mai et al. 1998]. In brief, to avoid impacting memory functionality the sense amplifier timing circuit and the wordline pulse generator circuit need to be redesigned. To avoid the redesign of these critical units and, moreover, not to increase bitline dynamic power dissipation we use *zig-zag share* circuit technique proposed in [Homayoun et al. 2008].

3.1 Zig-Zag Share Circuit

In this approach, sleep transistors are inserted in a zig-zag fashion [Min et al. 2003], [Horiguchi et al. 1993] keeping the R_{peq} of the first and third inverters and R_{neq} of the second and fourth inverters constant. This technique keeps the fall time of the circuit the same as in the baseline circuit with no leakage control. However, the rise time of the circuit is affected by the zig-zag scheme. In addition, using one sleep transistor per inverter logic increases the area for the zig-zag scheme.

To improve both leakage reduction and area-efficiency of the zig-zag scheme, [Homayoun et al. 2008] proposed using one set of sleep transistors shared between multiple stages of inverters which have similar logic behavior, such as stage 1 and 3 in a

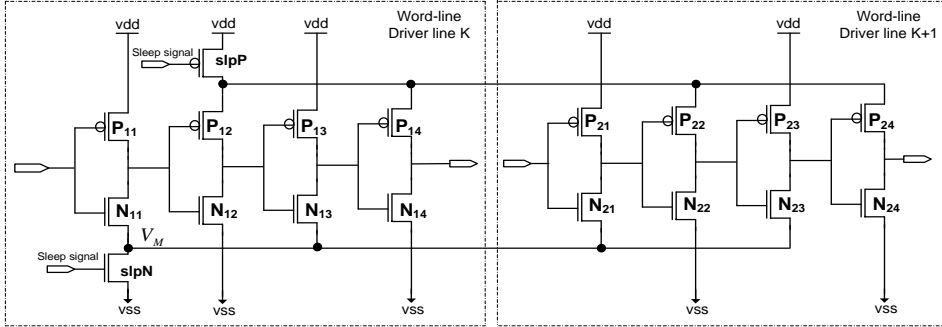


Figure 3. Zig-zag horizontal and vertical sharing circuit

studied chain of inverters. To further reduce leakage power, [Homayoun et al. 2008] proposed to also share one set of sleep transistors (slpN and slpP) vertically with adjacent rows of (wordline) drivers.

Figure 3 shows the *zig-zag horizontal and vertical sharing* circuit (in brief ZZ-HVS) when two adjacent wordline drivers share one set of sleep transistors. Intuitively, in vertical sharing (for instance for N_{11} and N_{21}), the virtual ground voltage (V_M in Figure 3) increases in comparison to when there is no vertical sharing.

Results in [Homayoun et al. 2008] show that using ZZ-HVS reduces the leakage power significantly, by 10 to 100X, when 1 to 10 wordlines share the same sleep transistors.

4. ZIGZAG-SHARE WITH MULTIPLE SLEEP MODES

As explained in Section 3, to maximize the benefit from the leakage savings of stacking sleep transistors we need to keep the bias voltage of the NMOS footer sleep transistor as low as possible (and for the PMOS header transistor as high as possible). The drawback of such biasing is its impact on wakeup latency and wakeup power of the circuit transitioning from sleep mode to active mode, which requires the voltage of virtual ground to reach the true ground. Such wakeup delay and its wakeup power overhead could significantly impact performance and diminish the power savings if incurred frequently. Appropriately sizing a sleep transistor (both footer and header) [Agarwal et al. 2006] and controlling its bias voltage are two effective ways to minimize the impact on wakeup delay and wakeup power overhead. For instance, increasing the gate voltage of the footer sleep transistor (in Figure 2(a)) reduces the virtual ground voltage (V_M), which subsequently reduces the circuit wakeup delay and wakeup power overhead. The negative impact of such biasing is a reduction in leakage power savings. To better explain this, let us examine the wakeup delay and wakeup power overhead as a function of sleep transistor size and its gate bias voltage.

The wakeup delay and power overhead are measured as the time and power required for the virtual ground voltage (V_M in Figure 2(a)) to be discharged through a sleep transistor to reach the ground level [Amrutur et al. 2000]. This wakeup delay is expressed as follows:

$$E_{wake up} = \frac{1}{2} C_{block} \times V_M^2 \quad \text{Equation 1}$$




Figure 4. A robust gate bias generator circuit.

$$T_{wakeUp} = \frac{V_M \times C_{block}}{I_{slp}} \quad \text{Equation 2}$$

where C_{block} is the total capacitance of the circuit block and I_{slp} is the current of the sleep transistor after it is turned on to wake up the block. Such wakeup overhead is decided by the equivalent load, as shown in EQ.1.

As shown in [Agarwal et al. 2006], the virtual ground voltage is linearly dependent on the sleep transistor gate voltage; increasing the gate voltage of the NMOS sleep transistor reduces V_M . According to EQ.1 and EQ.2, such reduction reduces the wakeup delay and wakeup power overhead.

Also as discussed in Section 3, increasing the gate voltage of the sleep transistor results in higher leakage power dissipation. In fact, by controlling the gate voltage of the footer and header transistors we can define different sleep modes where each mode has a different wakeup delay overhead and a different amount of leakage power reduction. We refer to this scheme as multiple sleep modes zig-zag horizontal and vertical sharing or in brief MZZ-HVS.

4.1 The Bias Generator

As described in the previous section, the benefit of multiple sleep modes is obtained by controlling the gate voltage of the sleep transistor. A stable and robust voltage reference generator is needed to ensure the performance of multiple sleep mode operation over process, voltage, and temperature variations.

A Conventional bandgap reference circuit consists of bipolar transistors and resistors that might occupy a noticeable amount of area. A CMOS voltage reference in Figure 4 consisting of transistors in saturation, triode, and subthreshold regions can provide sub-1-V reference required in our design. Transistors M1 – M4 in Figure 4 are always in saturation and generate supply independent current. Transistors M5 and M6 operate in the triode region as voltage controlled resistors and produce PTAT (Proportional to Absolute Temperature) voltage. Transistors M7 and M8 operate in the subthreshold region, behaving as bipolar transistors to produce CTAT (complementary to absolute temperature) voltage.




Figure 5. Reference voltage versus temperature in °C.




Figure 6. Wakeup delay and leakage power for different footer and header gate bias voltage pairs

The reference voltage is taken from the drain node of M6. Figure 5 shows the simulation results of reference voltage versus temperature from -40°C to 125°C , where only $\pm 1.6\%$ voltage variation around the nominal value is observed across this wide temperature range.

5. CIRCUIT EVALUATION

The proposed multiple sleep mode *zig-zag share* approach was applied to SRAM wordline drivers. A test experiment was set up in which the wordline inverter chain drives 256 one-bit memory cells. The memory cells and wordline drivers were laid out using Mentor Graphic IC-Station in a 65nm technology and simulated using Synopsis Hspice at typical corner (25 °) with extracted netlist and the supply voltage of 1.08V. Empirical results presented are for leakage power and wakeup delay for different pairs of sleep transistors' bias voltage.

Figure 6 shows normalized wakeup delay and normalized leakage power for different footer and header gate bias voltage pairs when ZZ-HVS is shared by 10 rows of wordline drivers. The figure shows a clear trade-off between the wakeup overhead and leakage power. Increasing the bias voltage of the footer transistor reduces the leakage power savings as well as the wakeup delay.

5.1 Impact of Sleep Transistor sizing on Leakage Power Reduction/Wakeup Delay and Propagation Delay

As discussed in Section 3, inserting sleep transistors in a zig-zag share fashion increases the circuit propagation delay as well as its rise time. Appropriately sizing the sleep transistor can minimize the impact on circuit propagation delay as well as on its rise time.

Assume that a $\beta\%$ increase in the wordline driver propagation delay is tolerable. Now let's find the appropriate size of sleep transistor so that the propagation delay does not increase beyond $\beta\%$. Delay of a gate without sleep transistor is expressed as:

$$T_d = \frac{C_L \times V_{DD}}{(V_{DD} - V_{il})^\alpha} \quad \text{Equation 3}$$

Delay of a gate with sleep transistor is expressed as,

$$T_{dsleep} = \frac{C_L \times V_{DD}}{(V_{DD} - V_x - V_{il})^\alpha} \quad \text{Equation 4}$$

C_L is the load capacitance at the gate's output, V_{il} is the threshold voltage of the inverter chain.

With a $\beta\%$ delay overhead allowed during active operation of the word line driver,

$$\frac{T_d}{T_{dsleep}} = (1 - \frac{\beta}{100}) \quad \text{Equation 5}$$

$$V_x = (\sqrt[\alpha]{(1 - \frac{\beta}{100})} - 1)(V_{DD} - V_{il}) \quad \text{Equation 6}$$

$$I_{sleep} = \mu_n \times C_{ox} \left(\frac{W}{L} \right)_{sleep} \left((V_{DD} - V_{th}) \times V_x - \frac{V_x^2}{2} \right) \quad \text{Equation 7}$$

For the word line driver, V_{il} is the same as V_{th} of sleep transistor.

$$I_{sleep} = \mu_n \times C_{ox} \left(\frac{W}{L} \right)_{sleep} \left(\left(\sqrt[\alpha]{(1 - \frac{\beta}{100})} - 1 \right) - \frac{\left(\sqrt[\alpha]{(1 - \frac{\beta}{100})} - 1 \right)^2}{2} \right) (V_{DD} - V_{th})^2 \quad \text{Equation 8}$$




Figure 7. Impact of sleep transistor sizing on propagation delay

- H. Homayoun, A. Veidenbaum, A. Sasan and Dean M. Tullsen

Table 1. Impact of sleep transistor sharing and sizing on the wakeup delay

#shared inverter chains	W(1X) (ns)	W(2X) (ns)	W(3X) (ns)	W(4X) (ns)	W(5X) (ns)	W(6X) (ns)	W(7X) (ns)	W(8X) (ns)
1	0.256	0.137	0.093	0.064	0.045	0.037	0.032	0.029
2	0.620	0.367	0.273	0.205	0.155	0.136	0.124	0.115
3	1.190	0.732	0.583	0.464	0.381	0.345	0.321	0.309
4	1.655	1.072	0.877	0.736	0.637	0.596	0.564	0.556
5	2.130	1.438	1.214	1.065	0.952	0.905	0.884	0.882
6	2.595	1.817	1.609	1.453	1.336	1.298	1.277	1.275
7	3.050	2.196	1.983	1.830	1.739	1.708	1.699	1.696
8	3.525	2.609	2.432	2.291	2.203	2.178	2.171	2.170
9	4.010	3.036	2.887	2.767	2.695	2.675	2.667	2.663
10	4.450	3.471	3.338	3.235	3.182	3.168	3.163	3.160

I_{sleep} is the maximum current flowing through the ground, which is the total discharge current of the word line driver. Assuming each inverter stage is sized a times larger than the previous stage the total discharge current is expressed as,

$$I_{discharge} = \frac{(V_{DD} - V_x)}{R_{eq,n}} + \frac{(V_{DD} - V_x)}{\frac{R_{eq,n}}{a^2}} + \frac{(V_{DD} - V_x)}{\frac{R_{eq,n}}{a^4}} + \dots = \frac{(V_{DD} - V_x)}{R_{eq,n}} \times (1 + a^2 + a^4 + \dots) \text{ Equation 9}$$




Figure 8. Relative leakage power reduction.

Table 2. Peripherals multiple sleep mode

power mode	wakeup delay (cycle)	leakage reduction (%)
basic-lp	1	42%
lp	2	75%
aggr-lp	3	81%
ultra-lp	4	90%

The size of sleep transistor can be found by substituting the discharge current into EQ.8.

$$\left(\frac{W}{L}\right)_{sleep} = \frac{I_{discharge}}{\mu_n \times C_{ox} \left[\left(\sqrt[3]{1 - \frac{\beta}{100}} - 1 \right) - \frac{\left(\sqrt[3]{1 - \frac{\beta}{100}} - 1 \right)^2}{2} \right] (V_{DD} - V_{th})^2}$$

The impact of sleep transistor sizing on propagation delay for ZZ-HVS circuit when 10 rows of wordline drivers share sleep transistor is shown in Figure 7. Increasing the size of sleep transistor reduces the propagation delay overhead, for instance increasing the size of sleep transistor by 8X reduces the impact on propagation delay from 10% down to near 1%.

Table 1 reports the impact of sleep transistor sharing and sizing on wakeup delay when each wordline drives 256 one-bit memory cells. Higher degree of sharing of the sleep transistor results in larger wakeup delay. Increasing the size of the sleep transistor reduces the wakeup delay. Increasing the number of shared wordlines reduces the benefit of using a larger sleep transistor on wakeup delay reduction.

Figure 8 reports the relative leakage power reduction as a function of sleep transistor size and the number of sharing inverter chains. These results show that the sleep transistor size has a small impact on leakage power savings.

Finally, note that the power overhead of waking up peripheral circuits from any low power mode is negligible, almost equivalent to the switching power of sleep transistors. Sharing a set of sleep transistors horizontally and vertically for multiple stages of (wordline) drivers makes the power overhead even smaller.

6. APPLYING ZZ-HVS TO L1 DATA AND INSTRUCTION CACHES

Some of the highest leakage units in the processor are the on-chip caches (such as L1 and L2). Thus, it makes a lot of sense to apply the ZZ-HVS circuit technique to these units. Note that we assume the small impact of ZZ-VHS on propagation delay (1%) can be tolerated and hidden in deep pipeline memories such as L1 cache and thus it does not degrade their operating frequency.

As explained above, there is a latency associated with waking up the SRAM peripheral circuitry. The overall time delay for transition to/from the standby mode, STL, is the sum of sleep transistors wakeup delay and a propagation delay of the sleep signal. Both of these delays increase as the memory area increases, especially the latter delay, because the sleep signal needs to be transmitted over a greater distance. Accordingly, depending on the memory size and configuration, there is a different wakeup delay overhead for a specific ZZ-HVS bias voltage. To find the STL delay for an SRAM array, SPICE and a modified version of CACTI were used to measure the wakeup delay of a sleep transistor and the propagation delay of the sleep signal, respectively. To estimate the propagation delay we assumed that the sleep signal has to be transmitted across the SRAM peripherals. Based on these experimental results different sleep modes were defined for DL1 and IL1 caches. Table 2 shows wakeup delay and relative peripheral circuit leakage reduction for these four different modes. The basic low power mode has the lowest leakage reduction but shortest wakeup delay. Next is lp mode which has higher leakage

savings. Aggressive and ultra sleep modes have even higher leakage savings but also a longer wakeup delay.

This section describes the architectural approach to control the ZZ-HVS sleep transistors in DL1 and IL1 caches² for two different types of embedded processors: a low-end and a high-end. We start by briefly describing the processor architectures and the experimental methodology used and then present results for different cache organizations.

6.1 Experimental Methodology

The approach proposed in this paper was evaluated for several processor architectures (shown in Table 3) which differ primarily in their memory hierarchy. A low-end processor uses 2, 4, 8 or 16KB L1 instruction and data caches and no L2 cache. The high-end processor has two levels of on-chip caches, with L1 cache size of up to 32KB. The rest of the processor architecture is similar to the ARM11 family of processors [ARM11 MPCore Processor].

Figure 9 shows the processor pipeline of ARM11, a single issue, out of order completion processor. It has two fetch stages and two stages for data cache access. The fetch stages fill a four-entry instruction fetch buffer [ARM11 MPCore Processor].

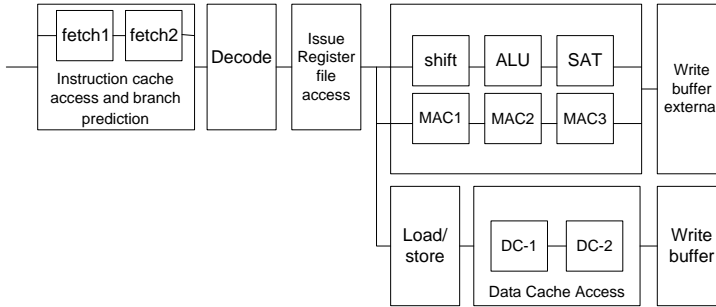


Figure 9. ARM11 processor pipeline

Table 3. Processor Architecture Configurations

	Low-end configuration	High-end
L1 I-cache	2-4-8-16KB, 4 way, 2	4-8-16-32KB, 8 way, 2 cycles
L1 D-cache	2-4-8-16KB, 4 way, 2	4-8-16-32KB, 8 way 2 cycles
L2-cache	none	64-128-256-512KB, 15 cycles
Fetch, dispatch	1 wide	1 wide
Issue	in-order, non blocking	in-order, non blocking
Memory	30 cycles	80 cycles
Instruction fetch queue	4	4
Load/store queue	4 entry	8 entry
Arithmetic units	1 integer, 1 f.p.	1 integer, 1 f.p.
Complex unit	1 integer, 1 f.p.	1 integer, 1 f.p. multiply/divide
Pipeline	8 stages	8 stages

² While it is possible to apply ZZ-HVS to other SRAM units such as the register file, reorder buffer, L2 cache, instruction cache, branch predictor, DTLB and ITLB, this paper studies the application in DL1 and IL1 cache only.

Clock speed	300 MHz	800 MHz
--------------------	---------	---------

The processor supports non-blocking and hit-under-miss operations in which it continues execution after a cache miss, as long as subsequent instructions are not dependent on cache miss data. The processor pipeline stalls only after three successive data cache misses.

An extensively modified MASE [SimpleScalar 4] simulator was used to model the architecture. The MiBench suite [MiBench Version 1.0] was used to represent a (low-end) embedded domain. SPEC2K benchmarks were used for the high-end processor. All benchmarks were compiled with the -O4 flag using the Compaq compiler targeting the Alpha 21264 processor. MiBench benchmarks executed for 500 Million instructions and SPEC2K for 500 Million instructions after fast-forwarding for 500 Millions instructions.

6.2 Reducing Leakage in L1 Data Cache

One approach to maximize leakage reduction in DL1 cache peripherals is to always keep them in ultra low power mode. However, this requires wakeup of DL1 peripheral circuits before each cache access and adds 4 cycles to the DL1 latency, which significantly reduces performance. One can put DL1 peripherals into the basic low power mode, which requires only one cycle to wakeup, and possibly even hide this wakeup latency, thus not degrading performance. However, this doesn't reduce leakage power significantly (see Table 2). To benefit from the large leakage reduction of ultra and aggressive low power modes while attaining the low performance impact of the basic-lp mode one has to dynamically change the peripheral circuit sleep modes. They should be kept in basic-lp mode during periods of frequent access and transitioned to aggr-lp or ultra-lp modes when accessed infrequently. The question is how to control such low-power mode transitions.

In this architecture it is known whether an instruction is a load or a store at least one cycle before cache access (during issue stage in Figure 9). As a result, DL1 peripheral circuits can be woken up one cycle prior to access. Thus there is no impact on DL1 access latency when its peripherals are in basic-lp mode. Similarly, one cycle of the wakeup delay can be hidden for all other low-power modes. Thus the effective wakeup delay of the DL1 cache is one cycle less than the delays shown in Table 2.

In addition, the following behavior of DL1 cache access has been observed and will be used to guide mode transitions. For low-end processors the DL1 is accessed infrequently, if at all, when there is one or more pending DL1 cache misses. Thus it can be put into ultra low power mode during miss service. For high-end processors the DL1 is accessed very infrequently while an L2 cache miss is being serviced. It is also access infrequently with no L2 cache misses but with multiple outstanding L1 misses and can be kept in aggressive low power modes.

Based on the above, the DL1 is kept in the basic-lp mode by default. The following state machines are proposed for controlling low-power mode transitions for different cache configurations.

6.2.1 Low-end architecture

For low-end architectures, the state machine shown in Figure 10(a) is proposed for controlling the DL1 low-power mode. Once a DL1 cache miss occurs the peripheral circuits transition from basic to lp mode. Given the miss service time of 30 cycles in this architecture, it is likely that the processor stalls during the miss service period. Occurrence of additional cache misses while one DL1 cache miss is already pending

further increases the chance of pipeline stall. Thus the DL1 peripherals are put into the aggressive low-power mode with more leakage savings.

Finally, the cache peripherals are put into the deepest low power mode, ultra-lp, once the processor stalls after DL1 misses occurs. A stall is detected in the issue stage and the processor put into ultra-lp when the processor does not issue any instructions for five consecutive cycles after a DL1 miss.

The processor returns to the basic-lp mode from any of the other low power states when one of the two following conditions are met:

- Stall condition removed; i.e instruction issue resumes
- All pending DL1 misses are serviced

Figure 11(a) reports the fraction of total execution time a 2KB DL1 cache spends in each of the power modes for the MiBench benchmarks. On average, 85% of the time DL1

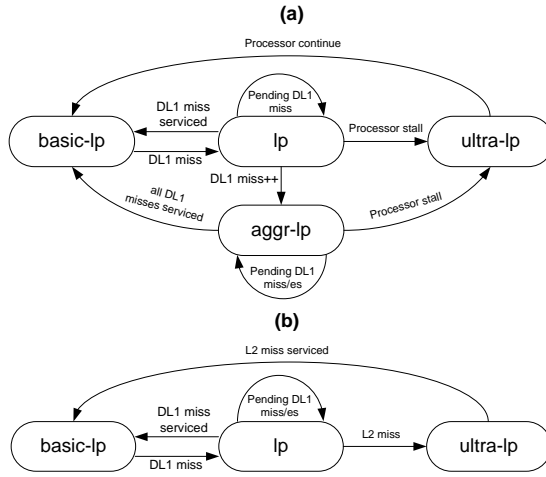


Figure 10. State machines to control DL1 cache peripherals in (a) low-end processor (b) high-end processor.




Figure 11. (a) Fraction of total execution time a 2KB DL1 cache spends in each of the power modes (b) Performance degradation of putting DL1 into low power mode (low-end architecture)




Figure 12. Fraction of total execution time DL1 is in different low power modes for (a) low-end and (b) high-end processors.

cache peripherals can be put into one of the low power modes. Most of the time is spent in the basic-lp mode, 58% of total execution time. 12%, 4% and 11% of total execution time is spent in lp aggr-lp and ultra-lp modes, respectively. In most benchmark ultra-lp mode does not have a noticeable contribution. Exceptions are rijndael, dijkstra and tiff2bw. In fact these are the benchmarks with the highest L1 miss rate (results not shown here) and as a result potentially large processor stall.

Figure 11(b) shows performance degradation for different DL1 cache sizes. The average performance degradation is less than 2% in all cases. Interestingly, the benchmarks which spend a considerable amount of time in one of lp, aggr-lp or ultra-lp modes have the most performance degradation, for instance dijkstra, lame and tiff2bw. This is understandable as transition to these modes incurs larger time delay.

Figure 12(a) shows the fraction of total execution time a DL1 spends in different low power modes for different cache sizes. Increasing the cache size reduces DL1 cache miss

rate and reduces opportunities to put the cache into more aggressive low power modes. This also reduces performance degradation for larger DL1 cache as can be seen in Figure 11(b).




Figure 13. (a) Fraction of total execution time a 4KB DL1 cache spends in each of the power modes (b) Performance degradation of putting DL1 into low power mode (high-end architecture)

6.2.2 High-end architecture

For high-end architecture we propose a simpler state machine to control the peripheral power mode (shown in Figure 10(b)). The major difference is that in this architecture the DL1 cache aggressively transitions to the ultra-lp mode right after an L2 miss occurs. Given a long L2 cache miss service time (80 cycles) the processor will stall waiting for memory. The cache returns to the basic-lp mode once the L2 miss is serviced.

Figure 13(a) shows the fraction of total execution time a 4KB DL1 cache spends in each power mode in the high-end processor. It shows that leakage power cannot be saved during only 10% of execution time. For the rest, the basic-lp mode has the highest contribution. Interestingly, in many benchmarks the ultra-lp mode has a considerable contribution, e.g. for ammp, applu, art and swim. In fact, these benchmarks have high L2 miss rate (not shown here), which triggers transition to ultra low power mode.

Figure 13(b) shows performance degradation for different DL1 cache sizes. The average performance degradation is less than 1.5%. Similar to the low-end architecture, here the benchmarks which spend considerable amount of time in one of lp, aggr-lp or ultra-lp mode have the highest performance degradation: ammp, applu, art and galgel.

Figure 12(b) shows the fraction of total execution time DL1 is put into different low power modes for different pairs of L1 and L2 cache sizes (X axis presents the L1-L2 cache size in the figure). Similar to the low-end architecture, increasing the cache size in any of high-end architecture reduces the relative power savings period.

6.3 Reducing Leakage in L1 Instruction Cache

The IL1 cache is accessed very frequently. The IL1 cache read port (data output drivers, read address input drivers, decoder/predecoder drivers) is being accessed more frequently than its write port (data input drivers, write address input driver and row decoder/predecoder drivers). Not that we assume the instruction cache has a separate read and write port (read port to read the instructions and write port to fill the instruction cache). The IL1 read port is accessed on average almost every cycle, while the write port is accessed very infrequently, once every 137 cycles in the Mibench benchmark (in low-end architecture) and once every 87 cycles in the SPEC2K benchmark (in high-end architecture).

Such different access pattern requires different control mechanism for reducing leakage. Our evaluation showed that making a write port one cycle slower has almost no impact on performance, but a one extra cycle of delay on every IL1 read (port) leads to noticeable performance degradation in some of the benchmarks. Therefore, the IL1 write port (data input drivers, write address input driver and row decoder/predecoder drivers) is always kept in basic-lp mode and is awakened on only when it is accessed. Both IL1 read and write ports (in both architectures) transition to the ultra-lp mode when the processor stalls. Our results indicate that this approach doesn't degrade performance.




Figure 14. (a) Fraction of total execution time a 2KB IL1 cache spends in each of the power modes (low-end architecture) (b) Fraction of total execution time a 4KB IL1 cache spends in each of the power modes (high-end architecture)

In Figure 14(a) we report the fraction of program execution time a 2KB IL1 cache spends in each of the power modes for MiBench benchmarks (in the low-end architecture). On average, 96% of the time the IL1 write port can be put into one of the low power modes. Most of the time is spent in the basic-lp mode, 82% of total execution time. 15% of total execution time is spent in ultra-lp mode. In most benchmark ultra-lp mode does not have a noticeable contribution. Exceptions are rijndael, dijkstra and tiff2bw. In fact these are

the benchmarks with the highest L1 miss rate (results not shown here) and as a result frequent processor stall. The IL1 write port transitions to high power mode (hp) for only 4% of execution time, on average. This is different than the read port where it spends most of the time in high power mode. In many benchmarks the IL1 read port is being accessed almost every cycle and as a result there is no opportunity to put its read peripherals into low power mode. Examples are bc, gsm and pgp. On average, the IL1 read port can be put into leakage savings ultra -low power mode for only 15% of total program execution time.

In Figure 14(b) we report the fraction of program execution time a 4KB IL1 cache spends in each of the power modes for the SPEC2K benchmarks (in high-end architecture). A similar trend is observed in this case where the write port can be put into a low power mode for a considerable portion of total program execution time while the read port needs to stay in high power mode most of the time. On average, the IL1 write port is in hp mode for 7% of total program execution time. This is 85% for the read port.




Figure 15. Fraction of total execution time IL1 read and write ports are in different low power modes for (a) low-end and (b) high-end processors.

Figure 15(a) shows the fraction of total execution time an IL1 cache spends in different low power modes for different cache sizes in the low-end architecture. Increasing the cache size reduces the IL1 cache miss rate and reduces opportunities to put the cache into more aggressive low power modes. Figure 15(b) present the results for high end architecture and for different pair of IL1 and L2 cache sizes (X axis present the IL1-L2 cache size in the figure). Similar to the low-end architecture, increasing the cache size in the high-end architecture reduces the relative power savings period.

7. POWER AND ENERGY-DELAY RESULTS

This section presents results for power reduction and energy-delay product. First, let us describe power and timing assumptions used. We used the relative leakage power reduction of various power modes reported in Table 2. Total dynamic power was computed as $N \cdot E_{\text{access}} / T_{\text{exec}}$, where N is the total number of accesses (obtained from simulation), E_{access} is the single access energy (from CACTI-5.1) and T_{exec} is the program execution time. Leakage power computations are similar, but leakage energy is dissipated on every cycle.

We also evaluate and included the power/area overhead of the controlling circuitry; a simple 2-bit state machine, a two 2-bit saturating counter for keeping the number of pending DL1 and IL1 misses and a 1-bit registers for keeping the L2 miss. Using Synopsys dc_compiler [Design Compiler] we synthesized the state machine with TSMC-65nm standard-cells which estimated the area overhead to be less than 200 gates (NAND2-gate). Using the estimated energy consumption reported by the dc_compiler and the activity of the state machine the overall power overhead estimated to be less than 0.6 mw (the power is almost 3nW/gate/MHz). Note that the exact area/energy

measurements require the detailed floorplaning and post placement and routing information.





Figure 16. (a) Leakage power reduction for the 2KB DL1 cache in the low-end architecture (b) Leakage power reduction for the 4KB DL1 cache in the high-end architecture

Figure 16(a) reports the leakage power reduction of individual benchmarks for the low-end architecture for the 2KB of DL1 cache. On average, DL1 leakage is reduced by 50%.

The fraction of leakage power reduction of each low power mode is also shown in the figure. Comparison of results in Figure 11 and Figure 16 (a) shows that while ultra-lp mode occurs much less frequently compared to the basic-lp mode, its leakage reduction is comparable to the basic-lp mode. The reason is that in the ultra-lp mode the peripheral leakage is reduced by 90%, almost twice that of the basic-lp mode.

Figure 16(b) shows the leakage reduction results for the 4KB DL1 cache in the high-end architecture. The average leakage reduction is almost 50%, with nearly half of it coming from the basic lp mode.






Figure 17. (a) Leakage power reduction for the 2KB IL1 cache in the low-end architecture (b) Leakage power reduction for the 4KB IL1 cache in the high-end architecture.

Figure 17 (a) reports the leakage power reduction of individual benchmarks for the low-end architecture for 2KB of IL1 cache. On average, IL1 leakage is reduced by 30%.

The fraction of leakage power reduction of each low power mode is also shown in the figure. Comparison of results in this figure and Figure 14(a) shows that while ultra-lp mode occurs much less frequently compared to the basic-lp mode, its leakage reduction is comparable to the basic-lp mode. The reason is that in the ultra-lp mode the peripheral leakage is reduced by 90%, almost twice that of the basic-lp mode. Another interesting observation from the figure is the small contribution of ultra low power mode in many benchmarks including apsi, eon, equake, mcf and mesa. In fact, in these benchmarks there is not much opportunity to put cache peripheral into ultra-low power mode as reported in figure 14.

Figure 17 (b) shows the leakage reduction results for the 4KB of IL1 cache in the high-end architecture. The average leakage reduction is almost 30%, with nearly tow third of it coming from the basic lp mode.




Figure 18. (a) Leakage power reduction (b) Total energy-delay reduction for DL1 cache.

Figure 18 shows the leakage power and energy-delay reduction of different DL1 cache sizes and processor configurations. On average, leakage is reduced by 42 to 52% for different configurations, with the bulk of the savings coming from the basic-lp mode. In both low-end and high-end architectures larger DL1 caches have lower leakage savings as their miss rates decrease. Overall, the energy-delay product reduction, unlike leakage power reduction, increases for larger cache sizes. The reason is that in smaller caches, the fraction of dynamic-energy per access to static energy is noticeably higher. As a result, for these small caches a large leakage reduction does not translate to large overall energy-delay reduction. This is different for larger caches, as their static power dissipation is

proportional to the dynamic power dissipation. The average energy-delay reduction varies from 9 to 21% for different architectures.




Figure 19. (a) Leakage power reduction (b) Total energy-delay reduction for IL1 cache.

Figure 19 shows the leakage power and energy-delay reduction of different IL1 cache sizes and processor configurations. On average, leakage is reduced by 24 to 33% for different configurations. In both low-end and high-end architectures larger IL1 caches have lower leakage savings as their miss rates decrease. Overall, the energy-delay product reduction, unlike leakage power reduction, increases for larger cache sizes. The reason is that in smaller caches, the fraction of dynamic-energy per access to static energy is noticeably higher. As a result, for these small caches a large leakage reduction does not translate to large overall energy-delay reduction. This is different for larger caches, as their static power dissipation is proportional to the dynamic power dissipation. The average energy-delay reduction varies from 5 to 11% for different architectures.

Applying the multiple sleep mode technique to both IL1 and DL1 simultaneously results in noticeable combined power savings. The overall processor power is reduced by up to 18.6% (an average of 11 to 14% for different architectures), assuming that the data cache and instruction cache consumes 16% and 27% of the total processor power, respectively [Montanaro et al. 1996].

8. CONCLUSIONS

This paper deals with an important problem of leakage energy in the DL1 and IL1 caches of embedded processors. While prior work has focused on reducing the power of the memory cell, this work focuses on peripheral SRAM circuits which account for 85% of overall leakage in 2KB to 16KB caches in 65nm technology. By defining multiple sleep modes and architectural control of sleep mode transitions, significant leakage energy reduction was achieved with no significant performance impact. In future process technologies of 32nm and below the peripheral circuit leakage is expected to be even higher and therefore the proposed approach will result in even higher energy savings.

REFERENCES

- Agarwal, A., Li, H., Roy, K., 2002, DRG-Cache: A data retention gated-ground cache for low Power, Design Automation Conference.
- Agarwal, K., Deogun, H., Sylvester, D. AND Nowka, K., 2006, Power gating with multiple sleep modes. In International Symposium on Quality Electronic Design.
- Amrutur, B. S., et al, 2000, Speed and power scaling of SRAMs, IEEE Journal of Solid State Circuits. vol. 35.
- Amrutur, B. S., et al., 2000, A replica technique for wordline and sense control in low-power SRAM's, IEEE Journal of Solid-State Circuits, vol. 33, No. 8.
- ARM11 MPCore Processor Revision: r1p0 Technical Reference Manual, infocenter.arm.com/ help/ topic/ com.arm.doc.ddi0360e/DDI0360E_arm11_mpcore_r1p0_trm.pdf.

- Calhoun, B. H. Honore, F. A., AND Chandrakasan, A., 2003, Design methodology for fine-grained leakage control in MTCMOS,” in Int. Symp. Low Power Electronics and Design, pp.104-109.
- Chun, S., AND Skadron, K., 2008, On-Demand Solution to Minimize I-Cache Leakage Energy with Maintaining Performance”, IEEE Transactions on Computers.
- Design Compiler, Synopsys Incorporation.
- Flautner, K., et al, 2002, Drowsy caches: simple techniques for reducing leakage power. IEEE ISCA.
- Gerosa, G., Curtis, S.D., Addeo, M., Jiang B., Kuttanna, B., Merchant, F., Patel, B., Taufique, M. AND Samarchi, H., 2008, A Sub-IW to 2W Low-Power IA Processor for Mobile Internet Devices and Ultra-Mobile PCs in 45nm Hi-K Metal Gate CMOS, in International Solid State Circuits Conference.
- Grossara, E., et al., 2006, Statistically Aware SRAM Memory Array Design, in International Symposium on Quality Electronic Design, ISQED.
- Homayoun, H., Makhzan M., AND Veidenbaum, A., 2008, Multiple Sleep Mode Leakage Control for Cache Peripheral Circuits in Embedded Processors, In Proceedings of the 2008 International Conference on Compilers, Architecture, and Synthesis for Embedded Systems, CASES, Atlanta, U.S.A.
- Homayoun, H., Makhzan, M., and Veidenbaum, A., 2008, ZZ-HVS: Zig-zag horizontal and vertical sleep transistor sharing to reduce leakage power in on chip SRAM peripheral circuits, in Proc. IEEE International Conference on Computer Design, pp.699–706.
- Horiguchi, M., Sakata, T., AND Itoh, K., 1993, Switched-source-impedance CMOS circuit for low-standby subthreshold current giga-scale LSI's, IEEE Journal of Solid-State Circuits.
- Kao, J., Narendra, S., AND Chandrakasan, A., 1998, MTCMOS hierarchical sizing based on mutual exclusive discharge patterns, Design Automation Conference.
- Kaxiras, S., et al, 2001, Cache decay: exploiting generational behavior to reduce cache leakage power. IEEE-ISCA.
- Khandelwal, V., AND Srivastava, A., 2004, Leakage control through fine-grained placement and sizing of sleep transistors, in Proc. ACM/IEEE Int. Conf. on Computer Aided Design, pp. 533-536.
- Kim, C. H., et al., 2005, A forward body-biased low-leakage SRAM cache: device, circuit and architecture considerations. IEEE Trans. on VLSI Systems, vol. 13.
- Kim, S. H. , AND Mooney, V. J., 2006, Sleepy keeper: a new approach to low-leakage power VLSI design, VLSI-SoC.
- Ku, J. C. , Ozdemir, S., Memik, G., AND Ismail, Y., 2006, Power Density Minimization for Highly-Associative Caches in Embedded Processors, GLSVLSI.
- Mai, K. W., Mori, T., Amrutur, B. S., Ho, R., Wilburn, B., Horowitz, M. A. , Fukushi, I., Izawa, T., AND Mitarai, S., 1998, Low-Power SRAM Design Using Half-Swing Pulse-Mode Techniques, IEEE Journal of Solid-State Circuits, vol. 33, no. 11, pp.1659-1671.
- Mamidipaka, M., Khouri, K.S., Dutt, N., AND Abadir, M. S., 2004, Analytical models for leakage power estimation of memory array structures. CODES+ISSS.
- Margala, M., 1999, Low power SRAM circuit design, In IEEE International Workshop on Memory Technology, Design, and Testing, 115–122.
- Meng, Y., Sherwood, T., AND Kastner, R., 2005, On the limits of leakage power reduction in caches. In International Symposium on High-Performance Computer Architecture.
- MiBench Version 1.0. <http://www.eecs.umich.edu/mibench/>.
- Min, K.-S., Kawaguchi, H. AND Sakurai, T., 2003, Zigzag super cut-off CMOS (ZSCCMOS) block activation with self-adaptive voltage level controller: an alternative to clock-gating scheme in leakage dominant era, IEEE International Solid State Circuits Conference.
- Montanaro, J., Witek, R. T., Anne, K., Black, A. J., Cooper, E. M., Dobberpuhl, D. W., Donohue, P. M., Eno, J., Hoepfner, W., Kruckemyer, D., Lee, T. H., Lin, P. C. M., Madden, L., Murray, D., Pearce, M. H., Santhanam, S., Snyder, K. J., Stephany, R., AND Thierauf, S.C., 1996, A 160-MHz, 32-b, 0.5-W CMOS RISC microprocessor, IEEE J. Solid-State Circuits, vol. 31, pp. 1703–1714.
- Mutoh, S., Shigematsu, S., Gotoh, Y., AND Konaka, S., 1999, Design method of MTCMOS power switch for lowvoltage high-speed LSIs,” in Proc. Asian and South Pacific Design Automation Conf, pp. 113-116.
- Nakagome Y., et al., 2003, Review and future prospects of low-voltage RAM circuits, IBM Journal of R&D.
- Nicolaescu, D., Salamat, B., Veidenbaum, A. AND Valero, M., 2006, Fast Speculative Address Generation and Way Caching for Reducing L1 Data Cache Energy. Proc. IEEE International Conference on Computer Design.
- Nii, K et al. 2004, A 90-nm low-power 32-kB embedded SRAM with gate leakage suppression circuit for mobile applications. IEEE Journal of Solid-state Circuits, Vol. 39, No. 4: 684-693.
- Nii, K., et al. 1998, A low power SRAM using auto-backgate-controlled MT-CMOS. In International Symposium on Low-Power Electronics, pp. 293-298.
- Park, J. C. AND Mooney III, V. J., 2006, Sleepy stack leakage reduction. IEEE Trans. VLSI Syst. 14(11): 1250-1263.

Reducing Leakage in Data Cache Peripheral Circuits of Embedded Processors

- Powell M., et al., 2000 Gated-Vdd: A Circuit Technique to Reduce Leakage in Deep-Submicron Cache Memories. IEEE-ISLPED.
- Rabaey, J. M., Chandrakasan, A. P. AND Nikolić, B., 2003, Digital integrated circuits: a design perspective, Prentice Hall, Second. Edition.
- Ramalingam, A., Zhang, B., Davgan, A., AND Pan, D., 2005, "Sleep Transistor Sizing Using Timing Criticality and Temporal Currents," Proc. ASP-DAC.
- Rusu, S., et al, 2007, A 65-nm Dual-Core Multithreaded Xeon® Processor With 16-MB L3 Cache, IEEE Journal Of Solid-State Circuits, VOL. 42.
- SimpleScalar4 tutorial, <http://www.simplescalar.com/tutorial.html>.
- Semiconductor Industries Association, 2005, International Technology Roadmap for Semiconductors, <http://www.itrs.net/>.
- Takeyama, Y. Otake, H., Hirabayashi, O., AND Kushida, K. N, 2006, A Low Leakage SRAM Macro with Replica Cell Biasing Scheme. IEEE Journal Of Solid- State Circuits, Vol. 41, No. 4.
- Thoziyoor, S., Muralimanoahar, N., Ahn, J.H., AND Jouppi, N.P., 2008, CACTI 5.1 Technical Report, HP Laboratories, Palo Alto.
- Zhang, W., AND Hu, J. S., 2002, Compiler-directed instruction cache leakage optimization. In Proc. In International Symposium on Microarchitecture.